

ADITHYA SHANKARAN

Data Engineer

+91 8780135961 | adikshan11@gmail.com | Surat, India

[linkedin.com/in/adithya-shankaran](https://www.linkedin.com/in/adithya-shankaran) | github.com/adikshan11 | adithya-shankaran.vercel.app

Summary

Data Engineer at **Xebia** with 2 years building **Spark**, **BigQuery** and **dbt** pipelines on **Google Cloud** for Levi Strauss & Co.: migrated **13 e-commerce sources** off Databricks, cut a top pipeline's run cost by **98%**, and ship reliability features in the shared **Vertex AI** framework. Builder of **imaarat.ai**, open-source AI underwriting for Indian property.

Experience

Xebia | Levi Strauss & Co. (client)

Aug 2024 – Present

Junior Consultant (Jul 2026 – Present) | Technical Trainee (Aug 2024 – Jun 2026)

- Shipped **per-task timeouts** with named alerts in the shared **Vertex AI** pipeline framework and rolled them out to the e-commerce pipelines, so a stuck task now stops and alerts instead of running for hours.
- Building one reusable **Dataproc** pipeline template in the shared framework to replace the versioned templates behind **600+ scheduled pipelines**, with a hard time limit and success read from the job's own final state.
- Cut a top-cost **Spark** pipeline's run cost by **98%** (\$51 to \$0.91 per run) and ended its repeated 4-hour timeouts by pruning joins and partitions; moved **29 pipelines** to right-sized clusters.
- Migrated **13 e-commerce data sources** (OMS, Hybris, Shopify, Salsify and more) from Databricks to **Spark on Dataproc** and **BigQuery**, validating every pipeline against the legacy output.
- Validated **7 Hybris** commerce pipelines against the new **SQL Server Hyperscale** source before cut-over, and repaired **1,455 missing records** in a **32 TB** dbt backfill.
- Designed **change data capture** for order data with **Datastream** into BigQuery to replace third-party replication; built **Dataplex** DQ scans and a CI gate that blocks dbt models without one.
- Centre of Excellence: built **RAG** document extraction on **Weaviate** and **ChromaDB**, and a Terraform linter on **Claude** (AWS Bedrock); selected for Xebia's **Forward Deployed Engineer** program.

Projects

imaarat.ai – AI Underwriting for Indian Property  

LangGraph, Gemini, RAG, MCP, A2A, FastAPI, React, Postgres, dbt

- Rules score each proposal against official seismic, flood and cyclone data for **19,312 PIN codes** while AI reads hand-filled forms in **26 languages**; ships with evals, tracing, reviewer sign-off and CI/CD.

NYC Taxi Lakehouse 

PySpark, Delta Lake, dbt, DuckDB, Airflow

- Delta Lake lakehouse over **9.5M trips** that quarantines **3.9%** bad records with 6 business rules, blocks bad loads with a 7-check quality gate, and serves dbt marts with **63 tests**.

Bookmark'd  

React, PostgreSQL, Redis, REST APIs, JWT, Docker

- Book discovery platform built as a four-member final-year project; I built search, reviews and **JWT** login in React over **PostgreSQL** and **Redis**.

Skills

Languages: Python, SQL, PySpark

Data Engineering: Apache Spark, dbt, Delta Lake, Databricks, Airflow, CDC, data modeling, data quality

Google Cloud: BigQuery, Dataproc, Vertex AI Pipelines, Dataplex, Datastream, Dataflow, Cloud Storage, Pub/Sub

AI: LLMs, RAG, LangGraph, MCP, A2A, Weaviate, ChromaDB, Claude Code, GitHub Copilot

Tools: Git, GitHub Actions, Docker, Terraform, PostgreSQL, Jira, Confluence

Education

Kalinga Institute of Industrial Technology

Sept 2020 – June 2024

B.Tech in Computer Science & Communication Engineering | CGPA: 8.32 / 10

Bhubaneswar, India

Certifications & Achievements

GitHub Copilot X-AI Practitioner+, Xebia (2026) | Mentor, Xebia Tech-AI-Thon | CodeChef rating 1614, 151 LeetCode problems | NSS Community Lead | Grade 8 Keyboard, Trinity College London